

# Creating Seamless 3D Maps Using Radiance Fields

Sai Tarun Sathyan ss4005@g.rit.edu  
Department of Computer Science  
Golisano College of Computing and Information Sciences  
Rochester Institute of Technology  
Rochester, NY 14623

**Abstract**— *Creating seamless 3D maps from 2D images is a challenging task. Traditional methods of creating 3D maps, such as photogrammetry, require a large number of images to be taken from different angles. This can be time-consuming and expensive, and it can also be difficult to stitch the images together seamlessly. NeRF[1] (Neural Radiance Fields) is a recent development in computer vision that can be used to create 3D maps from a collection of 2D images, similarly Gaussian splatting[4] is also a recent method of creating 3d models from 2d images. NeRF works by learning a mapping from a 5D input (3D position, viewing direction) to a 3D output (color, density). This mapping is represented by a neural network, which is trained on the images. Once the neural network is trained, it can be used to render new views of the scene from any angle. In this project, I will explore the use of NeRF & Gaussian splatting to create seamless 3D maps of cities, towns, and college campuses. I will first collect a dataset of 2D images of the scene, then train a model on this dataset. Once the model is trained, I will use it to render new views of the scene from any angle. I will evaluate the quality of the 3D maps created by NeRF and Gaussian splatting to compare them to traditional 3D maps.*

**Keywords**—*Nerf, Radiance Field, Photogrammetry, Gaussian splatting, 3d maps, insert (key words)*

## I. INTRODUCTION

NeRF (Neural Radiance Fields) is a recent development in computer vision that can be used to create 3D maps from a collection of 2D images. NeRF works by learning a mapping from a 5D input (3D position, viewing direction) to a 3D output (color, density). This mapping is represented by a neural network, which is trained on the images [1]. Once the neural network is trained, it can be used to render new views of the scene from any angle.

NeRF has several advantages over traditional methods of creating 3D maps. In this project, I will explore the use of NeRF to create seamless 3D maps of cities, towns, and college campuses. I will first collect a dataset of 2D images of the scene, then train a NeRF model on this dataset. Once the NeRF model is trained, I will use it to render new views of the scene from any angle. I will evaluate the quality of the 3D maps created by NeRF and compare them to traditional 3D maps. I believe that NeRF has the potential to revolutionize the way 3D maps are created. NeRF is a powerful tool that can be used to create realistic and seamless 3D maps from a small number of images. NeRF provides a way to store 3D scenes in a compressed format (as NN weights) with fast query times (I-NGP) and multi-scale photorealistic renders. Gaussian splatting is a rasterization technique for rendering radiance fields in real time. It is different from previous methods in that it

represents the radiance field as a set of 3D Gaussians, which makes it more efficient to render and optimize. Gaussian splatting works by first converting the radiance field into a set of 3D Gaussians. Each Gaussian represents a small point in the scene, and the intensity of the Gaussian represents the radiance at that point. The Gaussians are then splatted onto the screen, taking into account their visibility and the angle of the camera [4]. Gaussian splatting is different from previous methods in that it is more efficient to render and optimize. This is because Gaussians are a simple and well-understood mathematical object. They can be rendered quickly using graphics hardware, and they can be optimized efficiently to reduce the number of Gaussians that need to be rendered. Both NeRFs and Gaussians have the potential to make 3D maps more accessible and affordable, and to open up new possibilities for applications such as virtual reality and augmented reality. My goal is to create a robust workflow and scripts to automate this workflow so that creating seamless 3d maps becomes easier and simpler for others who want to attempt the same.

## II. RELATED WORK

A. "NeRF: Representing Scenes as Neural Radiance Fields" by Mildenhall et al. (2020)

NeRF simplifies 3D reconstruction by modeling scene geometry and appearance within a neural network framework to generate high-quality 3D scenes from 2D images[1]. Despite the original NeRF paper laying the foundation for this novel approach, it did not address real-time constraints, scalability issues, or model optimization potential. This paper conceptualized and developed NeRF, demonstrating its ability to synthesize novel views of a scene from a collection of 2D images, revolutionizing 3D reconstruction. It serves as the fundamental basis upon which subsequent research builds, introducing the concept of neural radiance fields and outlining the mathematical framework for scene representation. Later papers, like "Instant NeRF" and "Zip NeRF," extend this work by addressing practical limitations and efficiency concerns while leveraging the foundational NeRF concept.

B. "Instant Neural Graphics Primitives with a Multiresolution Hash Encoding" by Thomas Muller et al. (2022)

Instant NGP addresses the real-time limitations of NeRF, enabling practical, interactive applications in augmented reality and robotics [2]. It optimizes NeRF's training and rendering processes, bridging the gap between NeRF's theoretical elegance and practical usability. While it does not fully explore scalability for very large scenes or further model optimization, it is a significant contribution that makes real-time 3D reconstruction achievable.

C. "Zip-NeRF: Anti-Aliased Grid-Based Neural Radiance Fields" by Jonathan T. Barron et al (2023)

Zip NeRF by Jonathan T. Barron et al. enhances the efficiency and speed of neural radiance fields, making them more scalable and efficient for large-scale scenes[3]. It builds upon NeRF by optimizing the inference process with novel techniques like hierarchical sampling and optimization strategies. Zip-NeRF improves upon the original NeRF paper by getting rid of aberrations in the 3d model and improving the anti-aliasing, depth of field capabilities of NeRF. While Zip NeRF does not delve deeply into addressing dynamic scenes or explore additional use cases beyond real-time 3D reconstruction, it is a significant contribution to the NeRF literature, making large-scale 3D reconstruction more feasible and practical.

D. 3D Gaussian Splatting for Real-Time Radiance Field Rendering by Bernhard et al (2023)

The paper introduces a new approach to radiance field rendering called 3D Gaussian splatting, which achieves real-time rendering of high-quality radiance fields, even for large and complex scenes [4]. 3D Gaussian splatting is more efficient than traditional radiance field representations because it represents the radiance field as a set of 3D Gaussians, which can be rendered using a simple ray-marching algorithm. This achievement is significant because it makes it possible to render realistic 3D scenes in real time for a wide range of applications, such as augmented reality, virtual reality, and special effects. Overall, the paper is a significant contribution to the field of radiance field rendering. It introduces a new approach that is more efficient and versatile than previous approaches. 3D Gaussian splatting has the potential to revolutionize a wide range of applications. This method significantly advances the field of novel-view synthesis and real-time display of 3D scenes.

### III. SYSTEM FRAMEWORK

Creating a 3d model from 2d images is possible using photogrammetry, neural radiance fields and gaussian splatting. I have decided to move forward with neural radiance fields and gaussian splatting techniques since these two methods have overtaken the classic photogrammetry technique. Photogrammetry fails to accurately capture the density of reflective surfaces and transparent objects. When trying to create 3d maps of cities in which most buildings contain reflective surfaces i.e windows it becomes a huge problem to use photogrammetry. The first three papers mentioned in the previous work all depend on neural radiance fields at its core. The final paper relies on gaussian splatting to create 3d models. Neural radiance fields work in a unique way, when it comes to creating 3d models from 2d images. A static scene is represented as a continuous 5D function that outputs the radiance emitted in each direction ( $\theta, \phi$ ) at each point ( $x, y, z$ ) in space, and a density at each point which acts like a differential opacity controlling how much radiance is accumulated by a ray passing through ( $x, y, z$ ) [1].

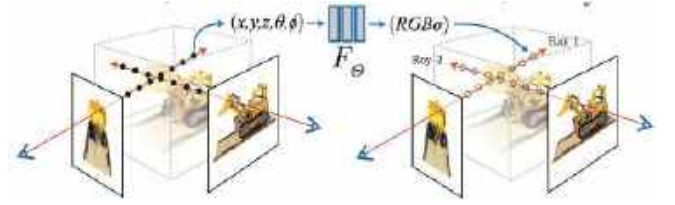


Fig 1: Neural Radiance Field Input Image Processing

As seen in Figure 1, NeRF represents a single 5D coordinate ( $x, y, z, \theta, \phi$ ) into a single volume density and view-dependent RGB color [1]. NeRF consists of two neural networks. The first one is called the coarse network. In 3d space this shoots out rays from each image and performs hierarchical volume sampling. It samples multiple coarse points spread out through that ray. Each coarse point is evaluated using the coarse network. The coarse network's job is to predict on which particular point on that ray is the volume high and on which point is the volume low. The coarse network rules out which points on this ray are empty. The second neural network used is called the fine network. This neural network focuses on the points at which the volume is greater than zero. It focuses on the non-empty region and tries to fine tune it. It focuses the computation on non-empty regions and fine tunes what it sees to make the volume density as accurate as possible.

The other method I will be using in this paper is gaussian splatting. Gaussian splatting allows real-time rendering for scenes captured with multiple photos, and creates the representations with optimization times as fast as the most efficient previous methods for typical real scenes[4].

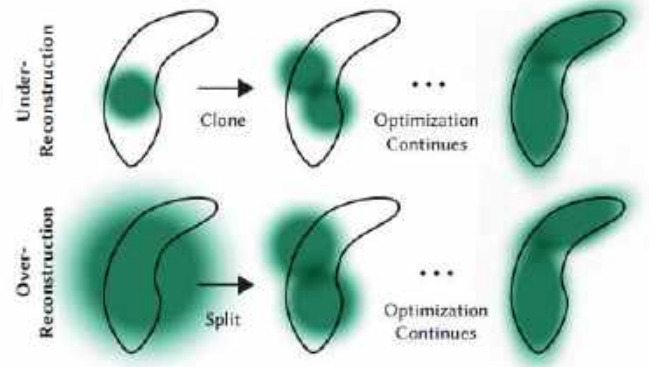


Fig 2: Gaussian Optimisation

Gaussian splatting starts by using the structure from motion algorithm to estimate a point cloud from the multiple pictures provided as input. Now instead of trying to create a 3d model from the polygons in the point cloud, it converts it into gaussians which take the shape of an ellipsoid. These gaussians now enter an iterative optimization process which fine tunes the gaussian to make it match the original images. In here adaptive density control takes place which either splits up the gaussians if it's too big or joins multiple gaussians if they are too small or transparent. Then the training starts using a process of rasterization, which projects the three dimensional gaussian back onto a two dimensional surface to accurately calculate the depth of each gaussian and the details of each gaussian are compared

to the original image and then gaussian gets fine tuned as shown in Fig 2.

Gaussian splatting, unlike photogrammetry or neural radiance fields, does not use ray tracing , path tracing or diffusion. Instead it relies on rasterization techniques for creating an accurate 3d model. This means it converts the underlying data directly into an image. Gaussian splatting requires millions of gaussians to function which requires several gigabytes of virtual ram.

These are the two main techniques which I will employ in order to create seamless 3d maps from 2d images. The "Instant Neural Graphics Primitives with a Multiresolution Hash Encoding" paper improves upon the original NeRF method by making it exponentially faster than the original one. Zip-NeRF is a different approach to NeRF where it projects cones instead of rays through the 2d image, the authors believe this improves depth perception. Unfortunately the code for this is unavailable. So for this project I will be using NGP and gaussian splatting.

#### IV. METHODOLOGY

In this section, I will describe the methodology I used to create 3D maps of the scene using NeRF and 3D Gaussian splatting. NeRF is a technique that uses a neural network to learn the relationship between 3D positions and the color and density of light at those positions. This allows NeRF to generate realistic views of the scene from any angle. 3D Gaussian splatting is a technique that uses a set of Gaussian functions to represent the 3D volume. It is typically used to render 3D volumes from a fixed viewpoint.

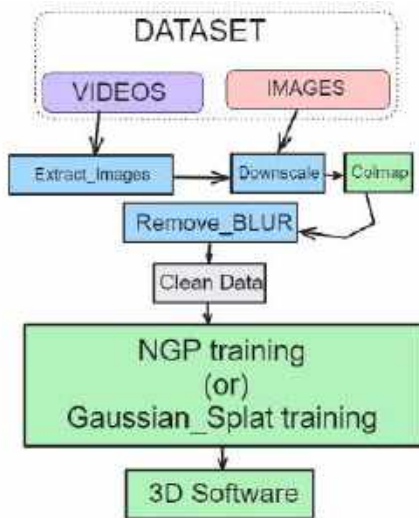


Fig 3: 3D model creation workflow

Even though the training methods and models are different , the method to create dataset for these two are the same. To capture the data, I used drone images of the scene and google satellite view as well. Then we move to data cleanup and pre-processing, next we feed this data into training. Once the NeRF neural network was trained and the 3D Gaussian splatting representation was generated, I used them to render realistic views of the scene from different angles. Finally, these 3d renders can be exported to a 3rd party 3D model viewing/editing software to manipulate it however we wish. I compared the rendered views from

NeRF and 3D Gaussian splatting to evaluate the two techniques.

##### A. Dataset Creation

The very first and the most important step of 3d model creation using either NeRF or gaussian splatting is the dataset capturing and creation technique. These models expect continuous images as raw input. The dataset can be created by capturing video or continuous images and then feeding it to Colmap script which tracks the camera positions for each image and lists it in a json [6].



Fig 4: Data capture techniques

As seen in Fig 4 There are various methods to capture different types of objects, some objects are small enough to be captured with the arc+boom technique which is going in a uniform circle around the object and completing a loop then moving upwards to repeat the same over and over. Whereas some objects are very large and need to be captured using truck+arc which involves going in a straight line and slowly arching around the corner of the object. Finally for objects that are really tall we can use a truck+pedestal technique which allows us to capture every detail laterally and then move upwards and repeat the same. All of these methods still require overlapping images to be taken, which is the most important part of dataset creation [5]. Overlapping images help to give Colmap script the context of where the camera is positioned and how it is moving around.



Fig 5: Image overlap technique

Fig 5 shows how images need to be captured in order to include overlapping photos in the dataset. This ensures that there are no duplications or multiple unknown artifacts in the final 3d model. Another important thing to avoid while capturing data is camera tilt. camera tilt will cause chromatic aberrations and blurry 3d models.

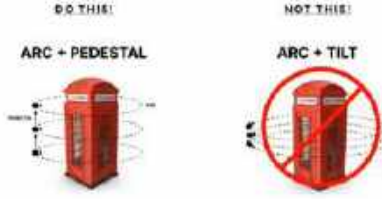


Fig 6: Avoiding camera tilt

If in any case you must capture a tall object you can use the arc+boom technique or arc+pedestal technique. Sudden changes in angles during dataset capture results in Colmap not being able to capture the camera movement, hence those images might not make it to the final dataset.

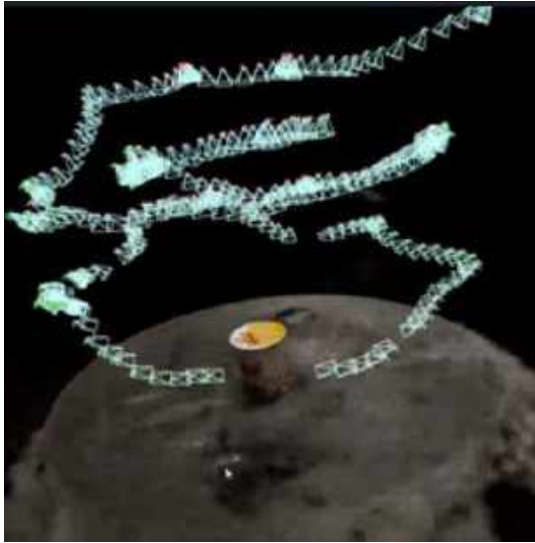


Fig 7: Avoiding camera tilt

Fig.7 shows a good real example of what camera movement should look like while capturing your images for the dataset. Following these techniques will lead to a better 3d model in the 3d, any mess up here will cost you a lot of wasted time and resources while training.

### B. Dataset Cleaning

After your dataset has been created you can directly input this data into training, if you have followed the dataset creation best practices you would get a decent 3d model. But there is still one way to improve this model and it is by removing blurry images out of the dataset. Blurry images obscures details in the objects and confuses NeRF and gaussian splatting models during training. because blurry images make the model hallucinate that there may be duplicate artifacts in this part of the object or it might lead the model to believe that this part of the object is transparent and the blur is caused by refraction. To avoid these issues we must remove blurry images. But manually reviewing every single image in your dataset and removing blurry images is time consuming and not a scalable solution. To overcome this obstacle I developed a blur removal script which does this for us. For the detection of blurry images Laplacian variance was used as the metric [7].

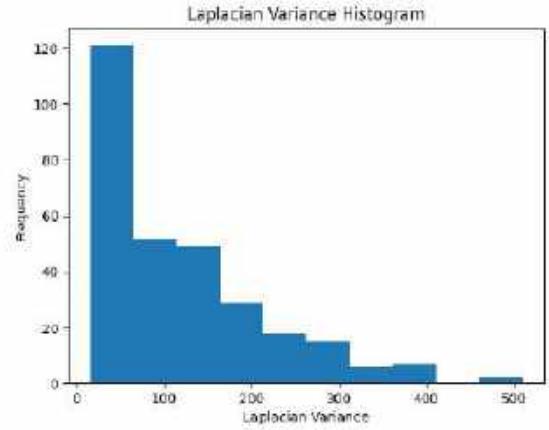


Fig 8: Dataset - Laplace variance distribution

Laplacian variance helps to calculate the sharpness of the image and this is done for every image inside our dataset. Next, the mean Laplacian variance is calculated for all images and the standard deviation is calculated as well. Using this I choose a threshold that is set to be anything lower than mean - standard deviation. Any image with a Laplacian variance lower than this set threshold is removed from the dataset. Laplacian variance is a measure of the sharpness of an image. It is calculated by computing the Laplacian of the image and then taking the variance of the Laplacian. The Laplacian of an image is a second-order derivative, which means that it highlights high-frequency features in the image.



Fig 9:3D model: blurry dataset vs non blurry dataset

As seen in Fig.9 NeRF can be sensitive to noisy data. If the dataset contains blurry images, NeRF will learn to represent the blurriness in the 3D model. This can result in a 3D model that is not as realistic as it could be. By cleaning the dataset using Laplacian variance, we can remove blurry images from the dataset. This will help NeRF to learn a more accurate representation of the scene, which will result in a more realistic 3D model.

### C. NeRF training

In NeRF (Neural Radiance Fields), the input to the neural networks comprises a 5D coordinate pair, consisting of 'x,' which represents the 3D spatial location, and 'd,' the 3D Cartesian unit vector denoting the viewing direction. These coordinates are pivotal in predicting two essential components for each input: volume density ( $\sigma$ ), which quantifies the amount of scene information at that location, and directional



emitted color (c), representing the color of light emitted from that 3D point in the specified viewing direction.

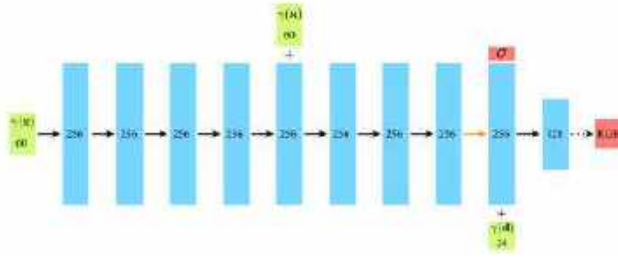


Fig 10: NeRF network architecture [1]

The training process involves two neural networks, namely the coarse and fine networks. The coarse network initially processes the 3D coordinates 'x' through a series of fully-connected layers with ReLU activations, employing 256 channels per layer. This operation serves to capture the volume density ( $\sigma$ ) and generates a 256-dimensional feature vector[1].

This feature vector is then combined with the camera ray's viewing direction 'd' and is further processed through an additional fully-connected layer with a ReLU activation and 128 channels, ultimately yielding the view-dependent RGB color (c). The fine network shares a similar architecture but conducts more network queries. It works to refine the representation established by the coarse network, thereby capturing finer details within the scene. In this process, it samples 192 points per ray, utilizing the previously obtained 5D coordinates. This fine-grained approach significantly enhances the quality of the rendered images. NeRF's training and optimization phase involve adjusting the weights of the MLP networks ( $\Theta$ ) to minimize the disparities between the predicted  $\sigma$  and c and the ground truth values extracted from the training images. The network parameters undergo fine-tuning throughout optimization, ensuring that the neural networks effectively model the scene's radiance field[1].

To maintain multiview consistency, NeRF employs an ingenious approach. It confines the prediction of volume density  $\sigma$  to be solely a function of the spatial location 'x,' while the prediction of RGB color 'c' depends on both the location and the viewing direction 'd.' This method equips the model to address non-Lambertian effects and effectively capture complex lighting phenomena, including specularities. In summation, NeRF's training process utilizes a pair of neural networks to develop a 5D representation of a scene based on multiple 2D images. These networks collectively consider the spatial structure and appearance of the scene. The coarse network captures the overarching scene structure, and the fine network refines intricate details, culminating in a detailed and view-consistent scene representation[1].

#### D. Gaussian Splatting training

The core of the training process revolves around optimizing the parameters of the 3D Gaussians. This includes refining the position, covariance,  $\alpha$ , and spherical harmonics (SH) coefficients representing color

for each Gaussian. This optimization is performed iteratively by rendering the scene and comparing the result to the training images to adjust Gaussian parameters. The training process needs to handle cases where geometry may be incorrectly placed due to 3D to 2D projection ambiguities, allowing the creation, destruction, or movement of geometry[4].

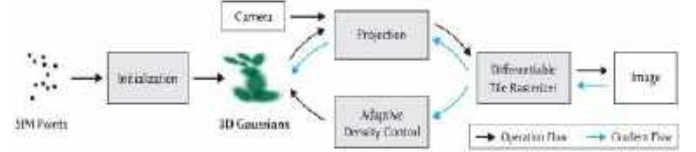


Fig 11: Gaussian splatting architecture [4]

The method adaptively controls the number and density of Gaussians over the unit volume to transition from an initial sparse set to a denser set that better represents the scene. This involves densification and removal of Gaussians that are nearly transparent ( $\alpha < \epsilon\alpha$ ) after an initial warm-up phase. Densification focuses on areas with missing geometric features and regions where Gaussians cover large areas in the scene.

Rasterization is a crucial part of the method, where Gaussian splats are rendered efficiently. A tile-based rasterizer is employed to pre-sort primitives for an entire image, which avoids the per-pixel sorting bottleneck. The method utilizes a sorting strategy based on GPU-accelerated radix sort to efficiently order the Gaussians and enable approximate  $\alpha$ -blending for pixel values. The pixel processing is parallelized, and each pixel accumulates color and  $\alpha$  values by traversing the lists of sorted Gaussians. The stopping criterion for rasterization is the saturation of  $\alpha$ , allowing for an arbitrary number of blended Gaussians to receive gradient updates, which is a notable advantage over previous approaches[4].

The optimization process for Gaussian parameters, along with adaptive control of Gaussians and rasterization, is based on Stochastic Gradient Descent (SGD) techniques, leveraging GPU acceleration for efficient processing. The loss function used combines L1 loss with a D-SSIM term [4]. The entire training process aims to optimize a scene representation that enables high-quality novel view synthesis, starting from sparse SfM points, and employs techniques to handle under-reconstruction and over-reconstruction situations.

It is a multi-step training process that involves adapting a set of 3D Gaussians to represent a scene effectively. It optimizes Gaussian parameters, controls their density adaptively, employs an efficient rasterization technique for rendering, and utilizes stochastic gradient descent for training. This comprehensive approach allows for the creation of a 3D model that can be used for high-quality novel view synthesis.

#### E. Exporting 3d model

This 3D model created by NeRF or Gaussian splatting could be exported into various 3D softwares. For NeRF the final 3D model created can be exported into blender or Meshlabs. The point cloud can be converted to a mesh to use for various things but the mesh is not perfect,

since the point cloud can have various empty spaces filling it up can result in an imperfect and jagged mesh.

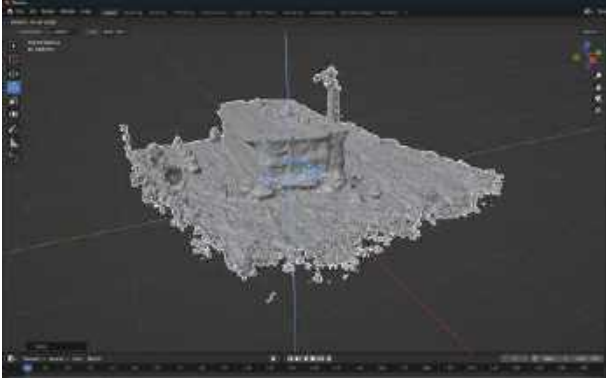


Fig 12: NeRF model in Blender

The problem with exporting it to blender as seen in Fig.12 is that it cannot reproduce the colors of the object created by NeRF.



Fig 13: NeRF model in MeshLabs

But meshlabs on the other hand is able to reproduce the original colors of the pointcloud. In the case of gaussian splatting we can use a pre-made unreal engine[9] plugin [8] to import the 3D model and play around with it.



Fig 14: Gaussian Splatt in Unreal Engine[8][9]

The good thing about this is that, since it's a game engine we can create an avatar to move around and interact with our 3D model (as seen in Fig.12) . The mesh object created by 3D gaussian splatting contains color and a cleaner mesh compared to the object produced by NeRF and this might be because gaussian splatting uses ellipsoids in 3D spaces to morph into the object it sees in the image and thus fills up the space in the point cloud, which makes the object have better

density distributions compared to the objects produced by NeRF. Gaussian splats can also be exported to unity for free.

## V. EXPERIMENTS

### A. Dataset Capturing Techniques

Given the set of rules mentioned in section 2.a I tried capturing various objects using different types of camera movements. These are the following data capture techniques I tried:

1. Keeping the camera parallel to the object and moving around the object while camera is focused on it
2. Moving around the object with the camera facing almost 90° away from it.
3. Moving around the object while the camera has an isometric view of the object

Technique 1 worked in some cases but failed most of the times to capture what was on top of the object or the top down view of it. It works only for very small objects. I tried employing technique 2 to capture buildings from a street view, the building was in the peripheral view of my camera as I walked around it. The end result was the 3D model being unable to map out the building properly in 3D space instead it created a 3D model of the building in a straight line rather than 4 sides of the building. Further testing can be done with 360° to get better results here. I lacked such equipment at the time of my experiments. Technique 3 proved to be the most useful method of data capture since it produced the most detailed 3D models as the end result. 3D models created using technique 3 as the data capture technique had the best 3D models in terms of details and accuracy.

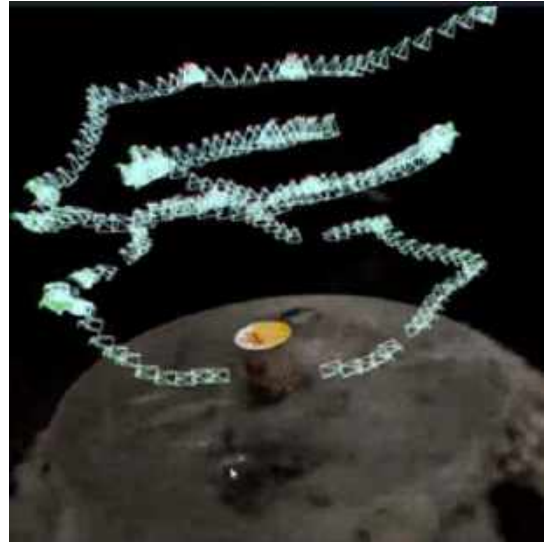


Fig 15: Camera movement for technique 3

As seen in Fig.15 The important thing to keep in mind is to never change the camera angle, I tried to maintain a 45° angle throughout my capture and moved in a circular motion around the object. Going around the object in a smooth motion, starting close to it and moving further away and higher up in the 2nd and 3rd rotation helps capture all details necessary for the 3d model recreation.



### B. Comparing NeRF and Gaussian Splatting

NeRFs and Gaussian splatting both have different 3D model recreation methods. In short, NeRFs produce a point cloud with colored points and Gaussian splatting produces 3D ellipsoid fields with color and transparency. When looking at these models in their respective GUI both the 3D models look almost the same but when we export these models into 3rd party 3D modeling software we get to see the true differences between the final results.

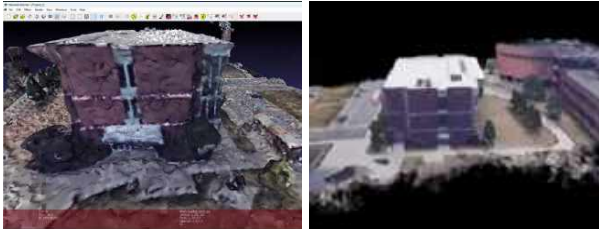


Fig 16: NeRF export compared with GUI preview

Here in Fig.16 we can see that the 3D model produced by NeRF looks fine in the NeRF GUI but when exported to meshlabs we lose a lot of detail and the model looks extremely choppy. This is because the point cloud itself is not perfect even i.e the point cloud density is not a perfect representation of what we see in real life. Due to parallax effect we may see it as a perfect 3D model in the NeRF GUI, but when exporting it to Meshlabs/Blender a thin Meshlayer is thrown over the object to maintain an even 3D mesh and here we get to see the true flaws in the 3D model created by NeRF.



Fig 17: Gaussian Splatting export

When Gaussian splatting gets exported to Unreal engine we get to see the model as seen in Fig.17

### C. 3D model comparison and scaling

After testing various single object scans I wanted to test how many objects can Gaussian splatting and NeRF capture accurately in one go. To find out the upper limit of these scans, I went to google satellite view[10] and captured multiple buildings in RIT campus and feed the training data to both the NeRF and Gaussian splatting model. These were the results of both the models outputs

starting from smaller single object scans to much wider multi object scans (same dataset was used to train both):



Fig 18: Ramenbox NeRF vs Gaussian Splatting



Fig 19: One Building NeRF vs Gaussian Splatting

The problem seen in Fig.19 is that the building is not being accurately reproduced in the gaussian splatting model. When I took a look at the training data to see what caused the problem, it turns out the dataset captured had slight camera tilt variation throughout the capture and this resulted in an extremely distorted 3D model while training Gaussian splats. This was an interesting insight to take note: Gaussian splatting is much more sensitive to camera tilts and this might cause the end result to be extremely distorted. After removing the images in the dataset that contained camera tilt I obtained this as the output for Gaussian splatting training:



Fig 20: One Building Gaussian Splatting model clean dataset

Next for the multi object scan these were the results:





Fig 21: Multi Building NeRF vs Gaussian Splatting

<b>NeRF</b>	176mb	8gb vram	90%	standard continuous images	1 min
<b>Gaussian Splatting</b>	256mb	24gb vram	95%	standard continuous images, but sensitive to camera shakes	5 min

Table 1: Gaussian Splatting vs NeRF (one building scan)



Fig 22: Multi Building NeRF vs Gaussian Splatting

In the multi-object scan test, Gaussian splatting has outperformed NeRFs render. When we look at the buildings from an isometric view in Fig.22 we can see that NeRF was not able to accurately reproduce the density of the point cloud in multiple locations, resulting in empty spaces in various locations. The Gaussian splatting model seems to have handled these areas a bit better, even though these areas contain rough and wide Gaussian patches, it is a better result than having an empty hole in the ground.

Comparison (for one building scan)	Memory efficiency	Resource consumption	Accuracy	Input data	Training speed

## VI. RESULTS

After seeing the performance of both NeRF and Gaussian splatting I decided to go forward with the Gaussian splatting model in order to create 3D maps. Here are some pics of the end results of the map creation:

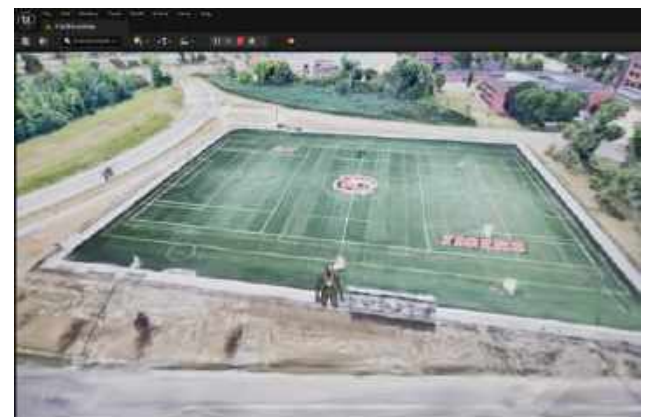
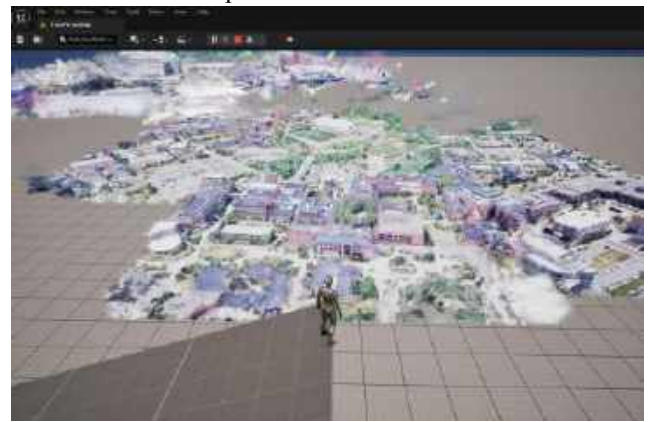


Fig 23: Map recreation in Unreal Engine



These assets in Fig.23 were created using google satellite view as input data and trained on the gaussian splatting model. Due to the drone flying restrictions on my university's campus I was unable to record the desired videos that I needed for training.

The exports from NeRF render looked subpar to the point cloud that was displayed on the preview window. This might be due to the NeRF model producing a very sparse point cloud which looks bad when a final mesh is applied on top of it, here is a real world object comparison with NeRF export and Gaussian splatting export:

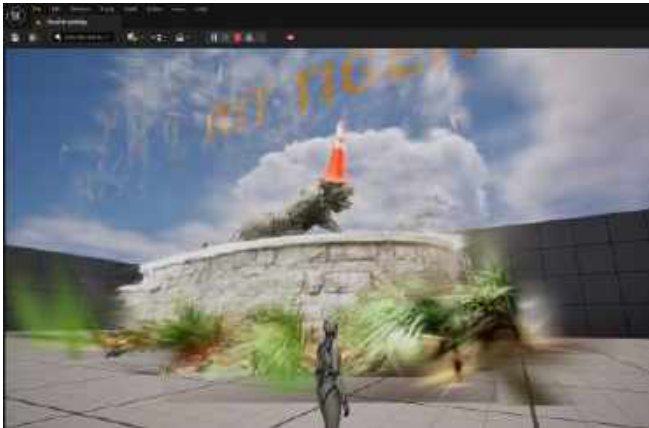


Fig 24: RIT tiger - Gaussian Splatting model



Fig 25: RIT tiger - NeRF model

## VII. FUTURE WORK

For future add ons to this project I would like to work on an unreal engine script to instantly take gaussian splats as grids and paste them into position when imported. This could make the workflow of creating a map much easier when it comes to large scale maps. I also would like to explore stable diffusion generative AI to amplify the looks of the final 3D model produced.

## REFERENCES

- [1] Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, Ren Ng, "NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis", Aug 2020. [Online] Available: <https://arxiv.org/abs/2003.08934>
- [2] Thomas Müller, Alex Evans, Christoph Schied, Alexander Keller, "Instant Neural Graphics Primitives with a Multiresolution Hash Encoding", May 2022. [Online] Available: <https://arxiv.org/abs/2201.05989>
- [3] Jonathan T. Barron, Ben Mildenhall, Dor Verbin, Pratul P. Srinivasan, Peter Hedman, "Zip-NeRF: Anti-Aliased Grid-Based Neural Radiance Fields", May 2023. [Online] Available: <https://arxiv.org/abs/2304.06706>
- [4] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, George Drettakis, "3D Gaussian Splatting for Real-Time Radiance Field Rendering", Aug 2023. [Online] Available: <https://arxiv.org/abs/2308.04079>
- [5] Jonathan Stephens, "NeRF dataset capture techniques Nvidia blog post", [Online] Available : <https://developer.nvidia.com/blog/getting-started-with-nvidia-instant-nerfs/>
- [6] COLMAP – Recover Camera Poses, [Online] Available : <https://liven.site/archives/2196>
- [7] Sagar, "Laplacian and its use in Blur Detection". [Online] Available : <https://medium.com/@sagardhungel/laplacian-and-its-use-in-blur-detection-fbac689f0f88>
- [8] "Unreal Engine 3D gaussian splatting plugin ". [Online] Available: <https://www.unrealengine.com/marketplace/en-US/product/3d-gaussians-plugin>
- [9] "Unreal Engine". [Online] Available: <https://www.unrealengine.com/en-US>
- [10] "Reconstruction of 3D building models from aerial images and maps". [Online] Available: <https://www.sciencedirect.com/science/article/abs/pii/S0924271603000583>